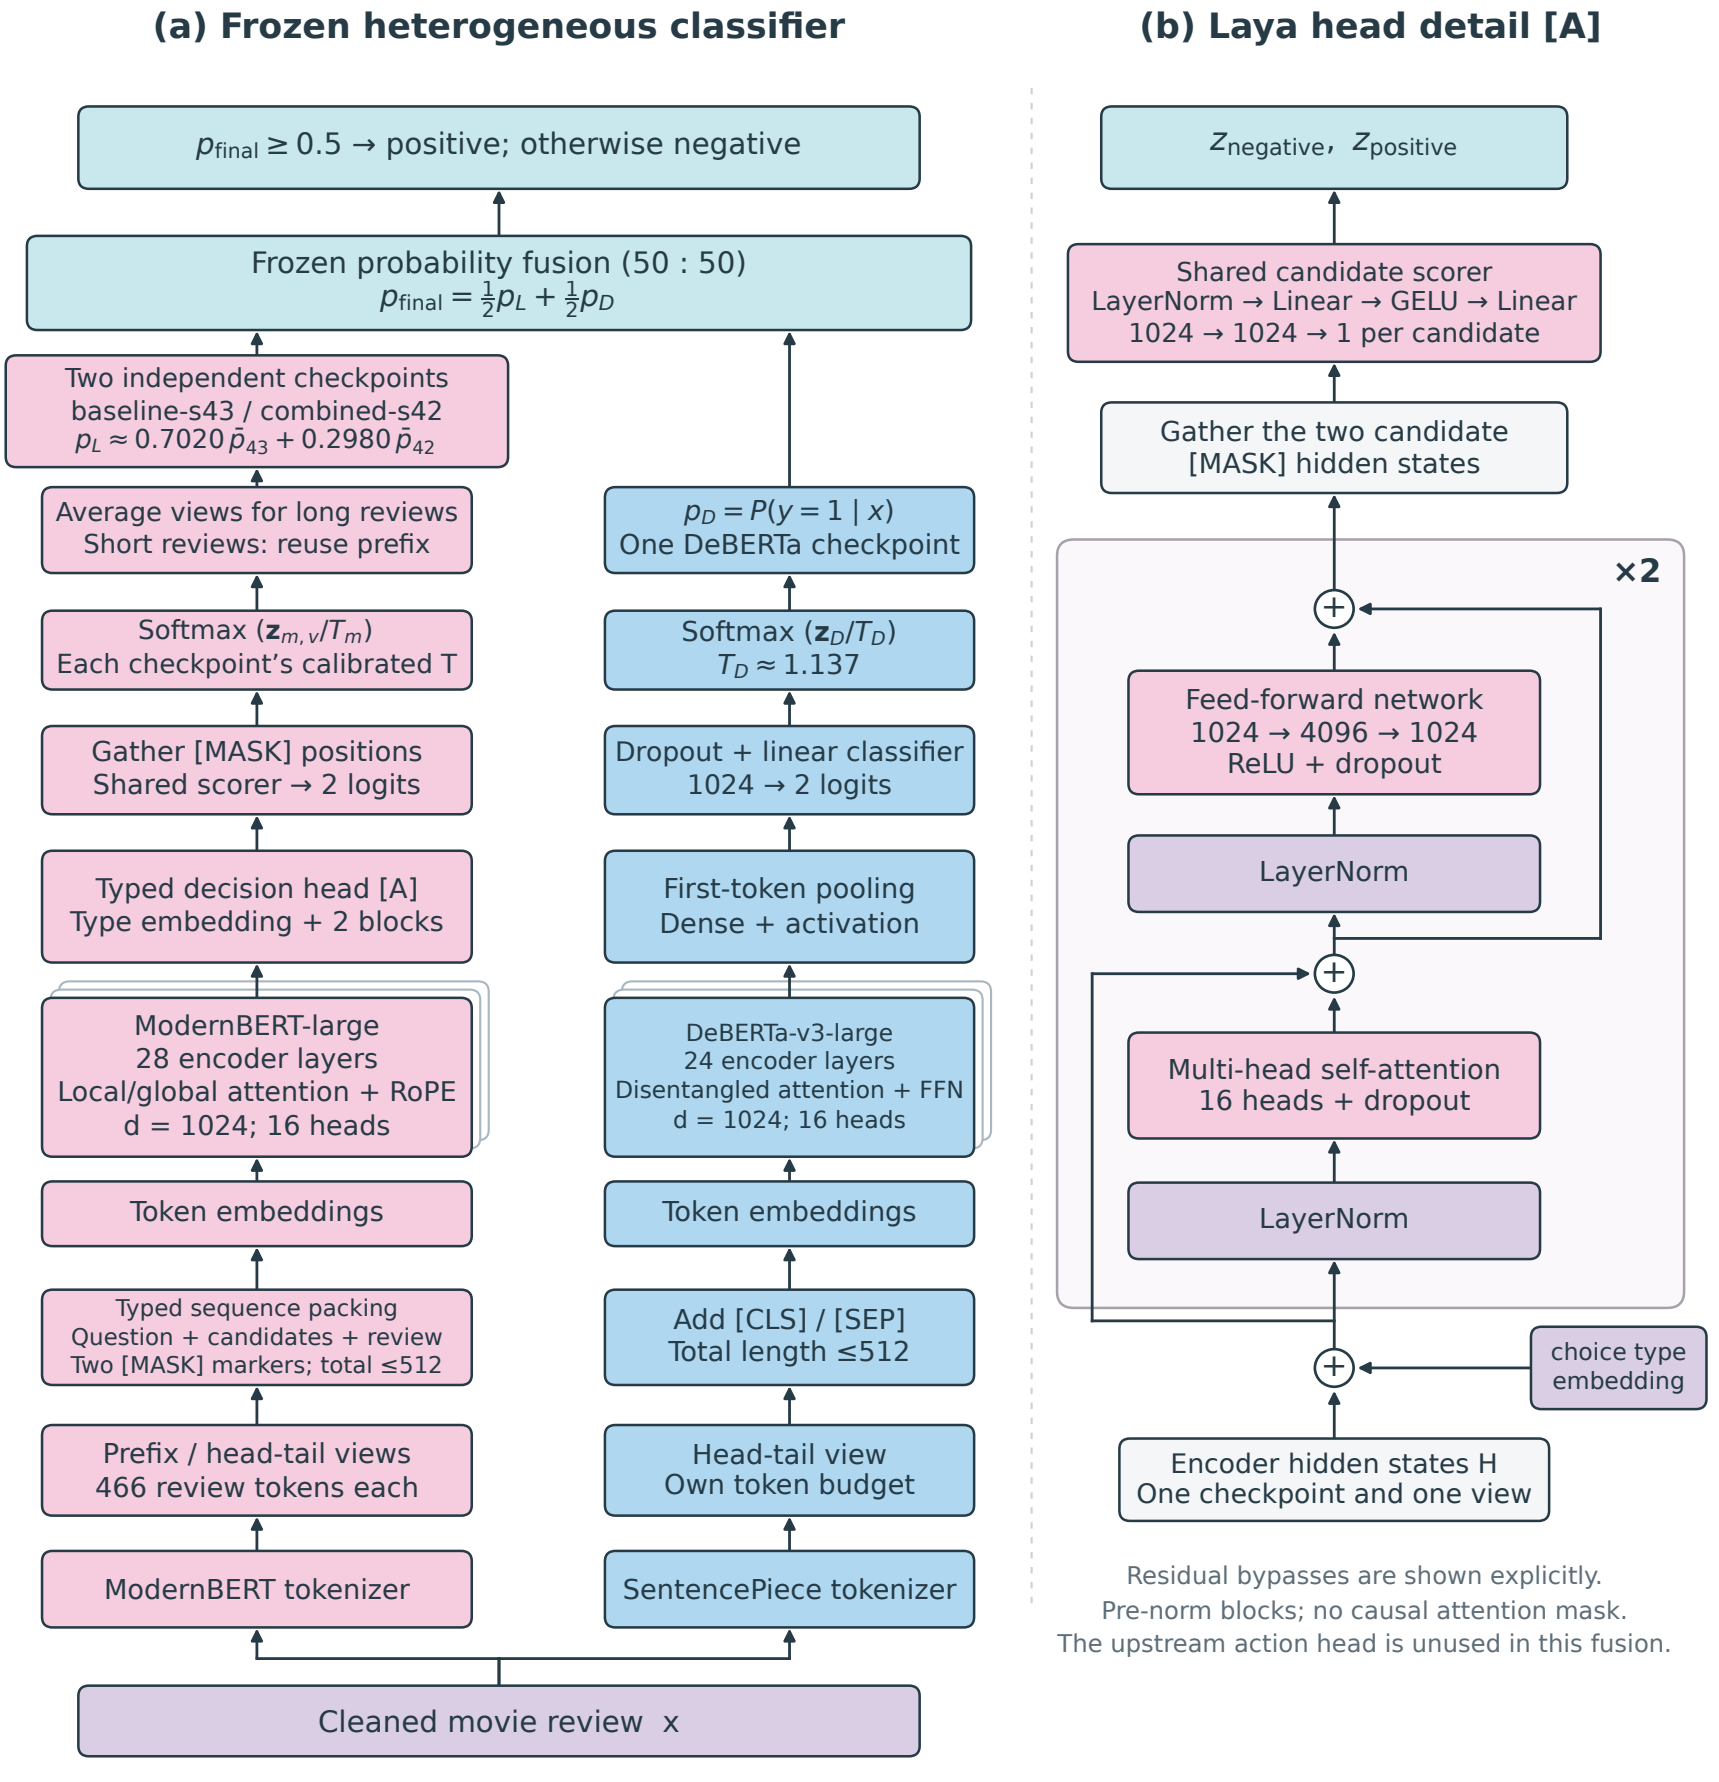
(a) Frozen heterogeneous classifier
$p_{\text{final}} \geq 0.5 \rightarrow$ positive; otherwise negative
Frozen probability fusion (50 : 50)
$p_{\text{final}} = \frac{1}{2}p_L + \frac{1}{2}p_D$
Two independent checkpoints
baseline-s43 / combined-s42
$p_L \approx 0.7020\,\bar{p}_{43} + 0.2980\,\bar{p}_{42}$
Average views for long reviews
Short reviews: reuse prefix
Softmax ($\mathbf{z}_{m,v}/T_m$)
Each checkpoint's calibrated T
Gather [MASK] positions
Shared scorer → 2 logits
Typed decision head [A]
Type embedding + 2 blocks
ModernBERT-large
28 encoder layers
Local/global attention + RoPE
d = 1024; 16 heads
Token embeddings
Typed sequence packing
Question + candidates + review
Two [MASK] markers; total ≤512
Prefix / head-tail views
466 review tokens each
ModernBERT tokenizer
$p_D = P(y = 1 \mid x)$
One DeBERTa checkpoint
Softmax ($\mathbf{z}_D/T_D$)
$T_D \approx 1.137$
Dropout + linear classifier
1024 → 2 logits
First-token pooling
Dense + activation
DeBERTa-v3-large
24 encoder layers
Disentangled attention + FFN
d = 1024; 16 heads
Token embeddings
Add [CLS] / [SEP]
Total length ≤512
Head-tail view
Own token budget
SentencePiece tokenizer
Cleaned movie review x
(b) Laya head detail [A]
$z_{\text{negative}}$, $z_{\text{positive}}$
Shared candidate scorer
LayerNorm → Linear → GELU → Linear
1024 → 1024 → 1 per candidate
Gather the two candidate
[MASK] hidden states
×2
Feed-forward network
1024 → 4096 → 1024
ReLU + dropout
LayerNorm
Multi-head self-attention
16 heads + dropout
LayerNorm
choice type
embedding
Encoder hidden states H
One checkpoint and one view
Residual bypasses are shown explicitly.
Pre-norm blocks; no causal attention mask.
The upstream action head is unused in this fusion.