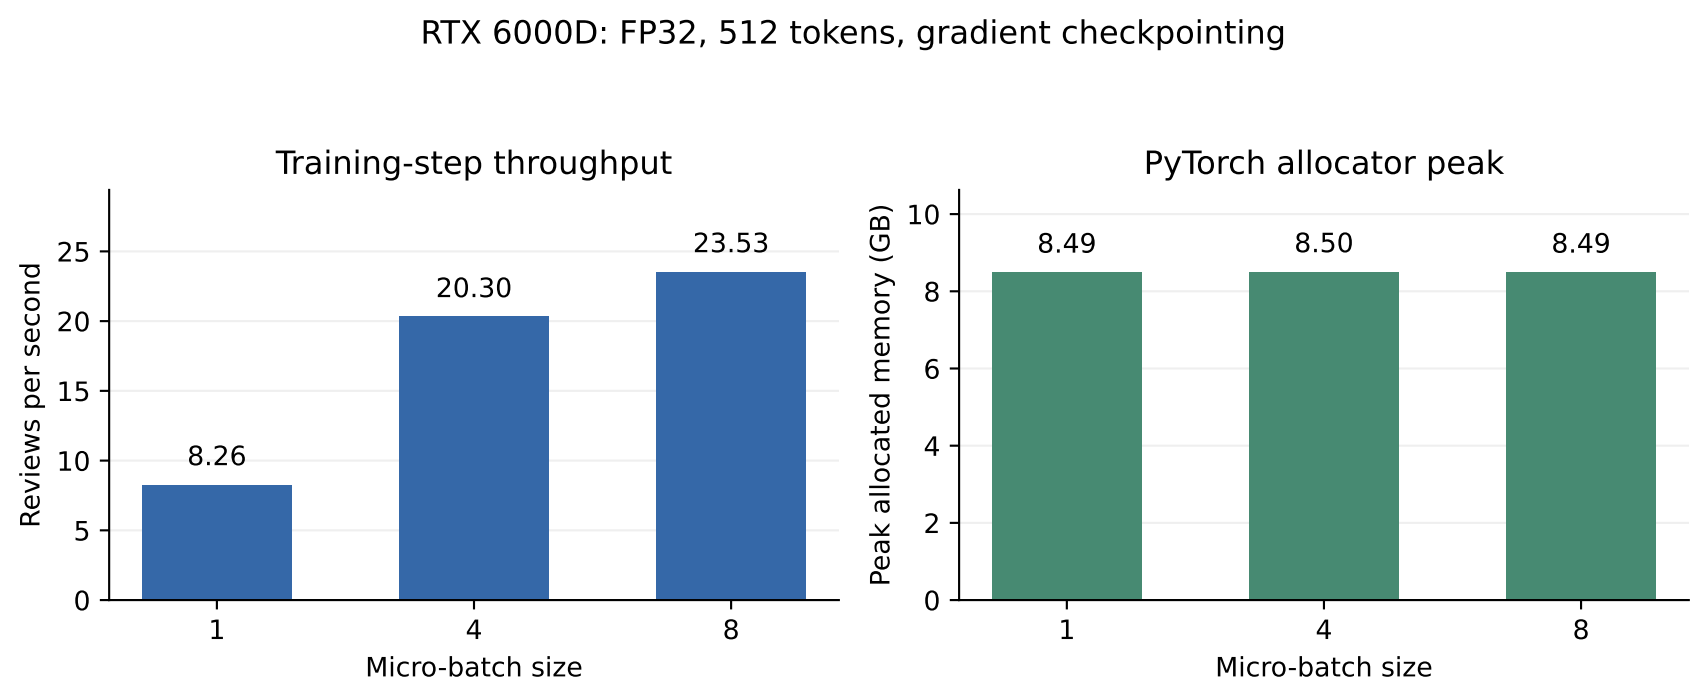
RTX 6000D: FP32, 512 tokens, gradient checkpointing
Training-step throughput
Reviews per second
0
5
10
15
20
25
8.26
20.30
23.53
1
4
8
Micro-batch size
PyTorch allocator peak
Peak allocated memory (GB)
0
2
4
6
8
10
8.49
8.50
8.49
1
4
8
Micro-batch size
2 warmup + 5 measured steps per batch; preloaded GPU inputs; zero LR.
Includes forward/backward/clip/AdamW; excludes tokenization, validation and checkpoint IO.